

Retrieval Augmented Generation (RAG)

Agenda

- Retrieval Augmented Augmentation (RAG) Quiz
- Building Blocks of RAG
- Chunking
- Semantic Search
- Evaluation

Let's begin the discussion by answering a few questions.

Which of the following represents the correct sequence of blocks for processing data in a retrieval-augmented generation (RAG) model?

A

Embedding -> Vector DB -> Data Chunk -> Retriever -> LLM

B

Data Chunk -> Embedding -> Vector DB -> Retriever -> LLM

C

Vector DB -> Data Chunk -> Retriever -> Embedding -> LLM

D

Data Chunk -> Embedding -> Retriever -> Vector DB

Which of the following represents the correct sequence of blocks for processing data in a retrieval-augmented generation (RAG) model?

A

Embedding -> Vector DB -> Data Chunk -> Retriever -> LLM

B

Data Chunk -> Embedding -> Vector DB -> Retriever -> LLM


C

Vector DB -> Data Chunk -> Retriever -> Embedding -> LLM

D

Data Chunk -> Embedding -> Retriever -> Vector DB

Retrieval Augmented Generation (RAG)



How does increasing the chunk overlap in a RAG system impact retrieval?

A

It reduces the number of retrieved chunks, improving efficiency

B

It increases redundancy, but helps maintain continuity in retrieved content

C

It eliminates the need for chunking altogether

D

It prevents the model from retrieving any irrelevant information

How does increasing the chunk overlap in a RAG system impact retrieval?

A

It reduces the number of retrieved chunks, improving efficiency

B

It increases redundancy, but helps maintain continuity in retrieved content

C

It eliminates the need for chunking altogether

D

It prevents the model from retrieving any irrelevant information

Chunking

The process of **breaking a document** (or text corpus) into **smaller, manageable pieces** (chunks) before storing them


Chunks are then **indexed and used to retrieve relevant information efficiently** when generating responses

512 => number of tokens in a chunk

16 => number of tokens overlapping between any two consecutive chunks

Increasing the chunk overlap parameter increases redundancy as there will be larger overlaps

But this helps maintain continuity in retrieved content as information might be spread across multiple chunks



Why do we need to create embeddings for our data in a Retrieval-Augmented Generation (RAG) system?

A

To improve keyword-based searching

B

To enable semantic similarity search

C

To reduce the amount of stored data

D

To make data more structured

Why do we need to create embeddings for our data in a Retrieval-Augmented Generation (RAG) system?

A

To improve keyword-based searching

B

To enable semantic similarity search


C

To reduce the amount of stored data

D

To make data more structured

Sentence Embeddings



Similar sentences are closer in the embedding space; dissimilar sentences are further apart

This enables **semantic search to retrieve sentences** (or chunks) from the data source that are **“similar” to the user query**

How does RAG enhance the capabilities of LLMs in natural language processing tasks ?

A

By limiting the access to external knowledge sources for generating responses

B

By reducing the need for fine-tuning LLMs on specific tasks

C

By integrating external knowledge sources to provide contextually rich and accurate responses

D

By reducing the model's size

How does RAG enhance the capabilities of LLMs in natural language processing tasks ?

A

By limiting the access to external knowledge sources for generating responses

B

By reducing the need for fine-tuning LLMs on specific tasks


C

By integrating external knowledge sources to provide contextually rich and accurate responses

D

By reducing the model's size

Advantage of RAG



In a Retrieval-Augmented Generation (RAG) system, which factor primarily ensures that the response correctly utilizes the provided information?

A

Clarity in conveying the task

B

Relevance of the response to the query

C

Faithfulness to the context

D

Length of the response

In a Retrieval-Augmented Generation (RAG) system, which factor primarily ensures that the response correctly utilizes the provided information?

A

Clarity in conveying the task

B


Relevance of the response to the query

C

Faithfulness to the context

D

Length of the response





Power Ahead!

